

Greenscape Pro -- AI Agent Strategy

Prepared for isthispossible.ai Take-Home Evaluation | Muhammad Aqib

Marcus's stated priority order (quoting -> follow-up -> crew coaching -> marketing) gets the top item right but the rest wrong. Crew coaching is real money but an order of magnitude below the real lever. Marketing is solving a problem he explicitly says doesn't exist: 'Quote [constrained]. I cannot keep up with the leads I have.' I ranked by dollar leverage and structural dependency, not by what was said out loud first.

1. Quote & Proposal Accelerator -- (P0 Build)

Purpose: Turn Marcus's raw site-walk notes into a structured, priced proposal draft in minutes, not days.

- Parses free-text site-walk notes against the existing pricing catalog (200+ line items).
- Auto-flags any proposal over \$30K for Carlos's render step (matches current rule).
- Produces a structured line-item draft Marcus reviews and approves, not sends blind.
- Routes approved proposals into the existing send flow (Slack notification as the trigger).

What it unblocks: Marcus is the bottleneck by his own admission -- 'I have to touch every proposal... nobody else knows how to do that.' This removes the interpretation step, not the judgment step.

Estimated ROI: 35-40% of qualified leads are lost to faster competitors during a 6-9 day quote cycle, on ~150 projects/year at \$28K average. Recovering half the lost deals by cutting cycle time to 1-2 days is worth \$1M+/year in recoverable revenue.

Why #1: Converting raw field notes into a priced scope is the exact structural constraint holding back revenue.

2. Post-Sign Milestone Tracker

Purpose: Automate the HOA/permit/deposit chase that currently stalls 8-12 projects at once.

- Tracks each signed project through HOA submission, permit status, and deposit payment.
- Sends automated reminder sequences to customers and internal nudges to Jenna.
- Escalates to a human only when a stage stalls past a defined threshold.

What it unblocks: Jenna's manual chasing, and the 2-3 week compounding delay that pushes crew scheduling.

Estimated ROI: 8-12 projects in limbo at \$28K average = \$224K-\$336K in delayed revenue at any given moment.

3. Build-Progress Update Agent

Purpose: Auto-generate Marcus-voiced progress updates from CompanyCam/Jobber milestones.

- Triggers on photo uploads or Jobber check-ins.
- Drafts a short, personal-sounding update Marcus can approve and send in one tap.
- Replaces the 'silence for 4-5 days -> anxious customer calls Jenna' pattern.

4. Closed-Lost Lead Reactivation Agent

Purpose: Personalized, Marcus-voiced re-engagement on the 1,400+ closed-lost lead pile.

- Pulls context from existing GHL notes per lead.
- Drafts individualized 'still thinking about your backyard?' style messages.
- Sends via GHL SMS/email with Marcus's approval.

Estimated ROI: Even a conservative 2% re-close on 1,400 leads at \$28K average is \$784K in latent revenue.

5. Lead Pre-Qualification Agent

Purpose: Filter obviously unqualified leads before they hit Marcus's calendar.

- SMS or voice bot asks 4-5 qualifying questions (budget, timeline, property ownership).
- Routes only qualified leads to Marcus's calendar queue.

System Dependency: Sits upstream of the whole operational chain: 5 -> 1 -> 2 -> 3, with 4 recycling unclosed leads back to the top.

Why #1 isn't simply the founder's stated #1

Marcus's stated priority is directionally right but operationally vague. 'Speed up quoting' could mean a dozen different builds. The data says the actual constraint is a single, specific interpretation step that only Marcus can currently perform -- that's the system worth automating.

Greenscape Pro -- AI Agent Strategy

Prepared for *isthispossible.ai* Take-Home Evaluation | Muhammad Aqib

What I excluded, and why

The marketing/content agent (Marcus's stated #4): He states outright twice that lead volume is not his constraint -- ROAS is healthy at 4-4.5x and he 'cannot keep up with the leads I have.' Building a marketing agent optimizes a non-bottleneck metric.

System Architecture Trade-offs & Scaling Bottlenecks

What breaks first at scale: Full pricing catalog in-context prompting works up to ~200 items. Beyond ~500 items, token costs rise and LLM attention drift degrades SKU matching accuracy.

Production solution: Two-stage Vector RAG pipeline (OpenAI text-embedding-3-small or pgvector in Supabase) to dynamically retrieve top 15 relevant items before generation.

What I'd do with more time: Direct GoHighLevel (GHL) SMS API webhooks and line-item confidence scoring.